



Dependency parsing

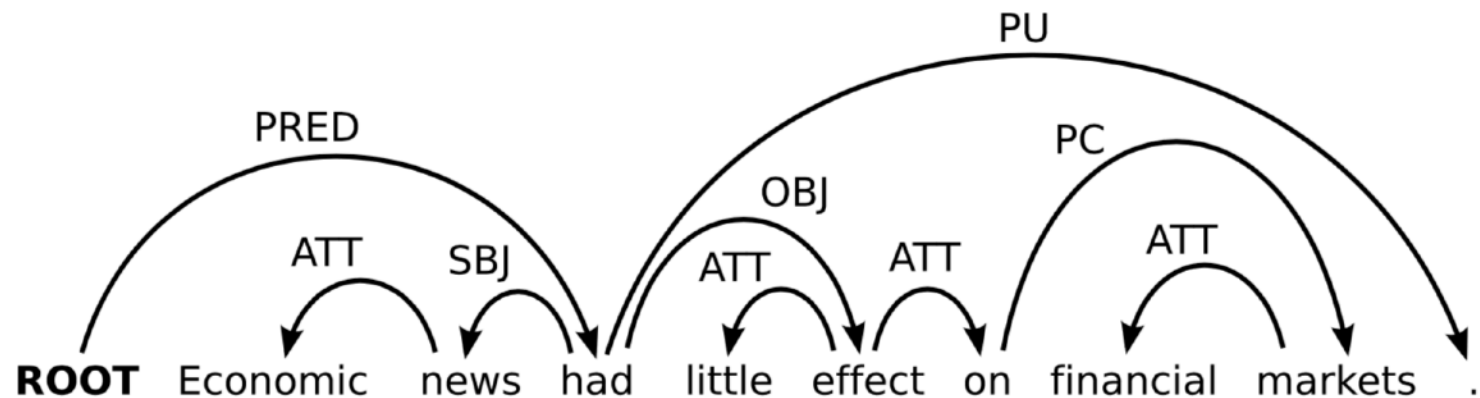
Benoît Sagot
Inria (ALMAAnaCH)

MVA — Speech and Language Processing — Class #6 — 24th February, 2020

Credit and disclaimer: some of the following slides are taken from, illustrated or inspired by presentations and article figures by Nivre, Dyer, Ballesteros, Kutuzov, Mooney, Rasooli and Tetreault

Introduction

- **Syntactic parsing of natural language**
 - Building the structure of natural language sentences
- **Dependency-based syntactic representations**
 - Long tradition in descriptive and theoretical linguistics
 - Have become popular in computational linguistics




Strategies for dependency parsing

- Graph-based parsing
- Transition-based parsing
- Other strategies

Graph-based parsing

- MSTParser (McDonald et al. 2005)
 - <http://www.seas.upenn.edu/~strctrln/MSTParser/MSTParser.html>
- Simplified version of the underlying idea:
 - Create all possible dependencies
 - Weigh them
 - Extract the optimal dependency tree
 - I.e. the tree that covers all words and minimises the overall weight of all retained dependencies

Arc-standard Transition-Based Parsing



Starting point

- **The basic idea:**

- Define a transition system for dependency parsing
- Learn a model for scoring possible transitions
- Parse by searching for the optimal transition sequence

- **Advantages:**


- Highly efficient parsing with low complexity
- Rich history-based feature models for disambiguation

- Cf. Nivre (et al.)

- <http://www.maltparser.org>

Formalising dependency trees

- A **dependency tree** is a **labelled directed tree** T with
 - a set V of nodes, labelled with wordforms (including the special “wordform” **ROOT**)
 - a set A of arcs, labelled with dependency types
 - a linear precedence order $<$ on V
- **Notation:**
 - Arc (w_i, l, w_j) connects head w_i to dependent w_j with label l
 - Node w_0 (labeled **ROOT**) is the unique root of the tree



Parser configurations

- A **parser configuration** is a triple $c = (S, Q, A)$, where
 - S = a stack $[\dots, w_i]_S$ of partially processed nodes,
 - Q = a queue $[w_j, \dots]_Q$ of remaining input nodes,
 - A = a set of labelled arcs (w_i, l, w_j) .
- **Initialisation:**
 $([w_0]_S, [w_1, \dots, w_n]_Q, \{\})$
(recall that $w_0 = \mathbf{ROOT}$)
- **Termination:** $([w_0]_S, [], A)$

Transitions for the “arc-standard algorithm”

- **Left-Arc(*l*)**

$$\frac{([\dots, w_i, w_j]_S, Q, A)}{([\dots, w_j]_S, Q, A \cup \{(w_j, l, w_i)\})} [j \neq 0]$$

- **Right-Arc(*l*)**

$$\frac{([\dots, w_i, w_j]_S, Q, A)}{([\dots, w_i]_S, Q, A \cup \{(w_i, l, w_j)\})}$$

- **Shift**

$$\frac{([\dots]_S, [w_i, \dots]_Q, A)}{([\dots, w_i]_S, [\dots]_Q, A)}$$

Example Transition Sequence

[ROOT]_S [Economic, news, had, little, effect, on, financial, markets, .]_Q

ROOT Economic news had little effect on financial markets .

Example Transition Sequence

[ROOT, Economic]_S [news, had, little, effect, on, financial, markets, .]_Q

action: Shift

ROOT Economic news had little effect on financial markets .

Example Transition Sequence

[ROOT, Economic, news]_S [had, little, effect, on, financial, markets, .]_Q


action: Shift

ROOT Economic news had little effect on financial markets .

Example Transition Sequence

[ROOT, *Economic*, news]_S [had, little, effect, on, financial, markets, .]_Q


action: Left-Arc(ATT)



Example Transition Sequence

[ROOT, news, had]_S [little, effect, on, financial, markets, .]_Q

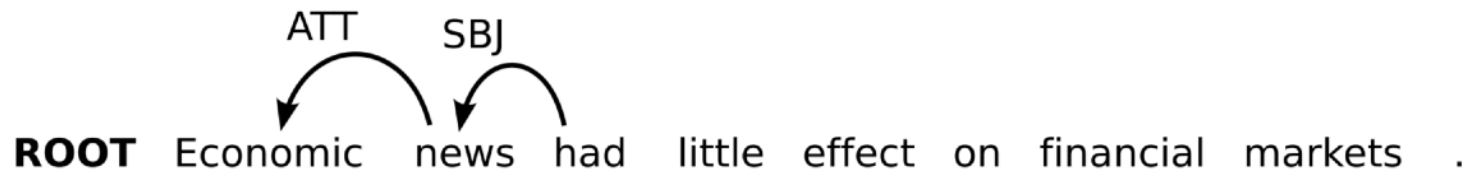
action: Shift



Example Transition Sequence

[ROOT, news, had]_s [little, effect, on, financial, markets, .]_Q

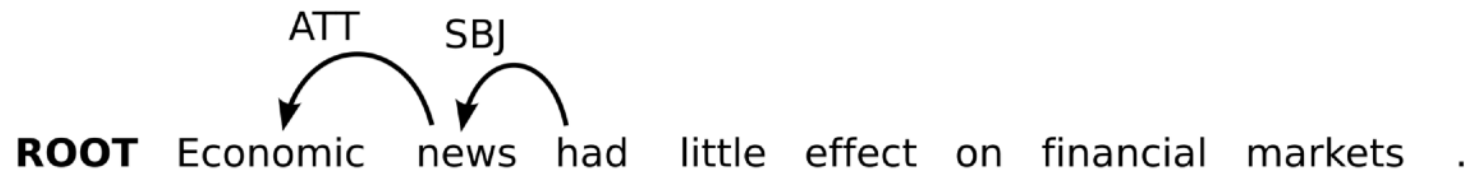
action: Left-Arc(SBJ)



Example Transition Sequence

[ROOT, had, little]_S [effect, on, financial, markets, .]_Q

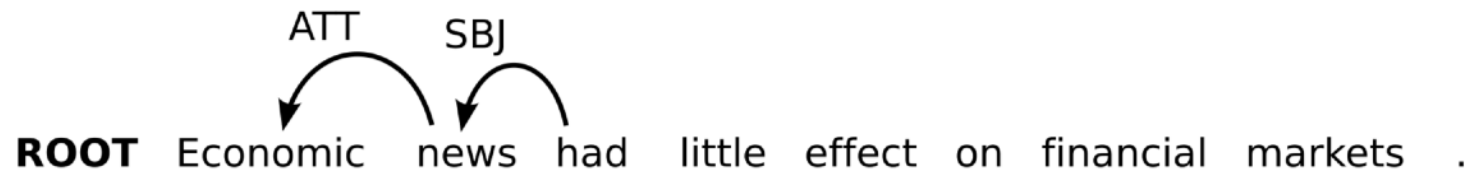
action: Shift



Example Transition Sequence

[ROOT, had, little, effect]_s [on, financial, markets, .]_Q


action: Shift



Example Transition Sequence

[ROOT, had, little, effect]_s [on, financial, markets, .]_Q


action: Left-Arc(ATT)



Example Transition Sequence

[ROOT, had, effect, on]_S [financial, markets, .]_Q


action: Shift



Example Transition Sequence

[ROOT, had, effect, on, financial]_s [markets, .]_q


action: Shift



Example Transition Sequence

[ROOT, had, effect, on, financial, markets]_s [.]_q


action: Shift



Example Transition Sequence

[ROOT, had, effect, on, financial, markets]_s [.]_q

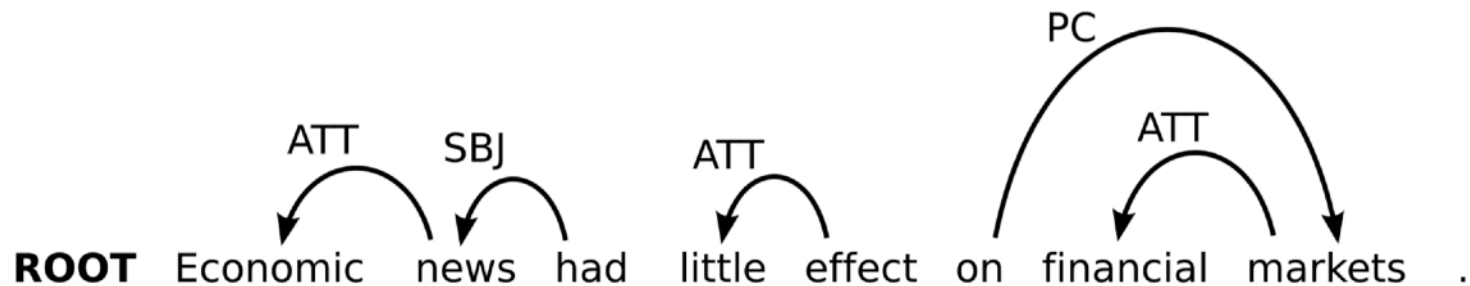
action: Left-Arc(ATT)



Example Transition Sequence

[ROOT, had, effect, on, markets]_s [.]_Q

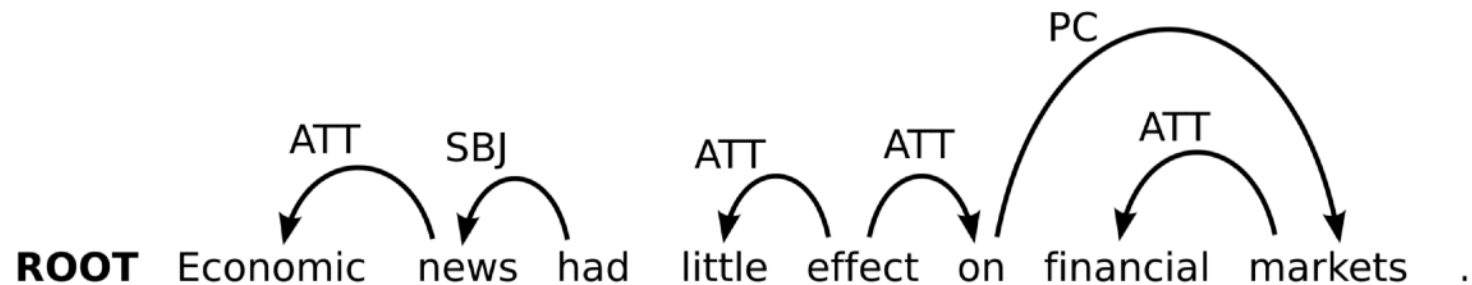
action: Right-Arc(PC)



Example Transition Sequence

[ROOT, had, effect, on]_s [.]_q

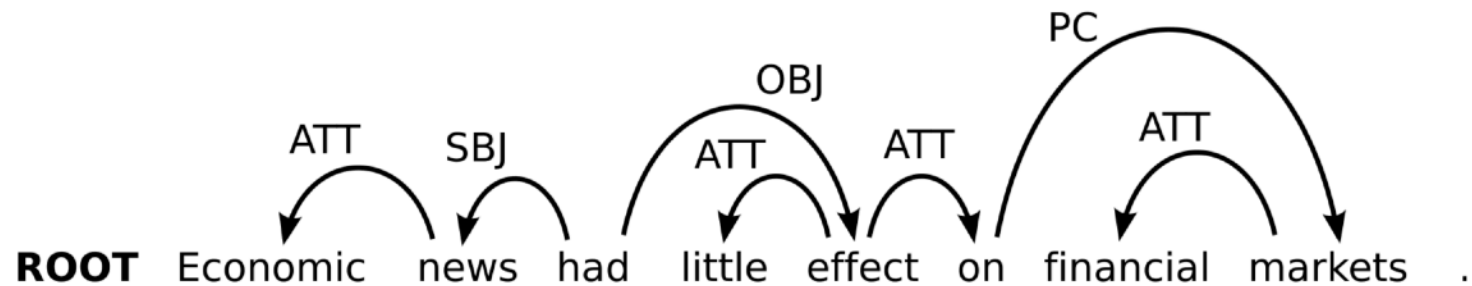
action: Right-Arc(ATT)



Example Transition Sequence

[ROOT, had, effect]_s [.]_q

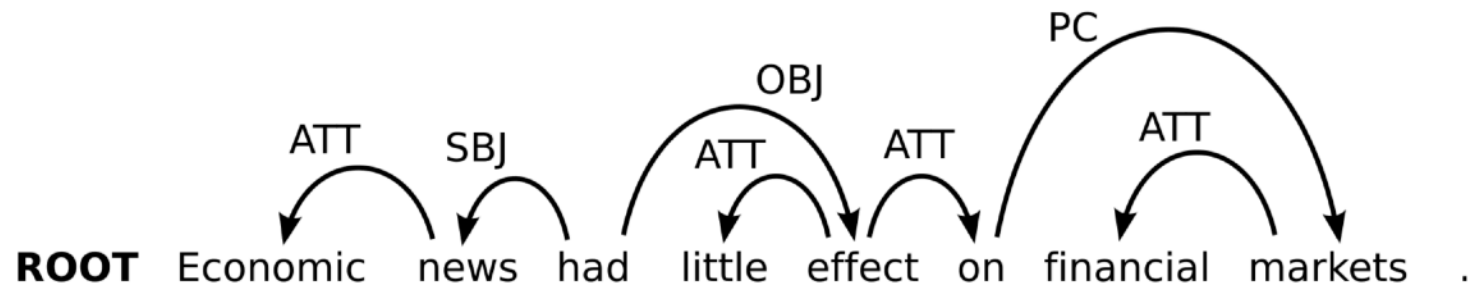
action: Right-Arc(OBJ)



Example Transition Sequence

[ROOT, had, .]s []q

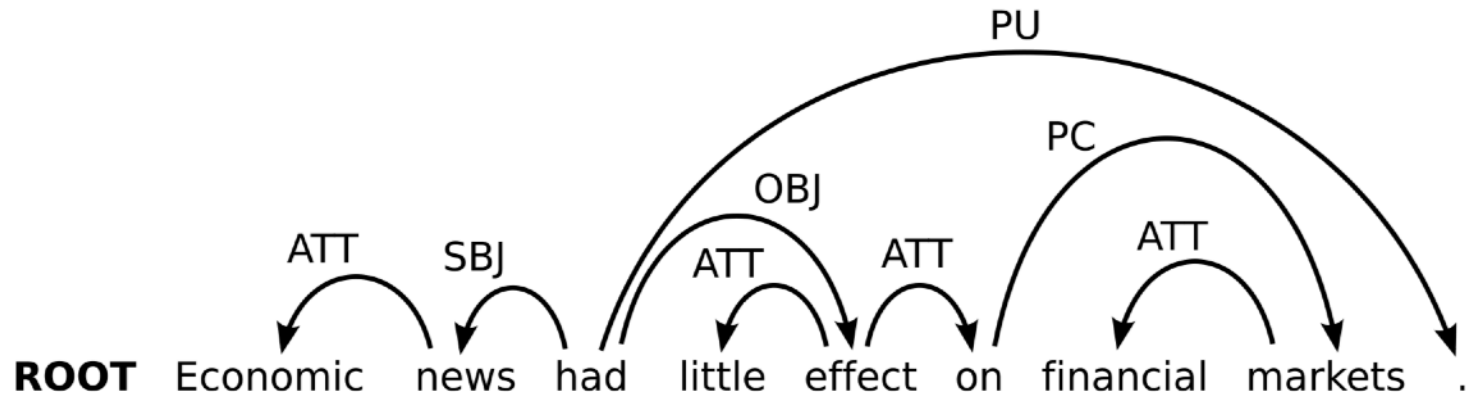
action: Shift



Example Transition Sequence

[ROOT, had, .]s []q

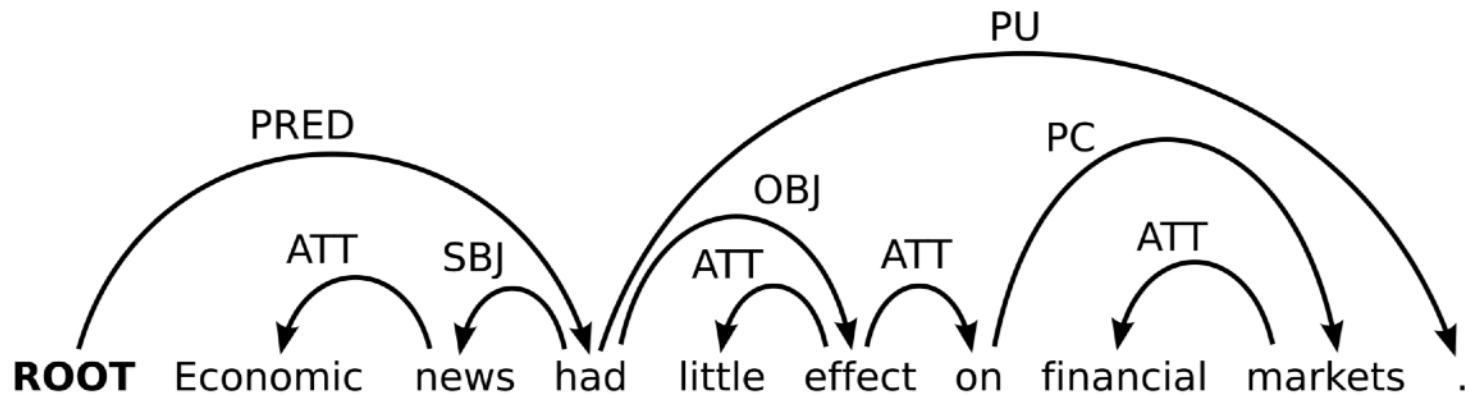
action: Right-Arc(PU)



Example Transition Sequence

[ROOT, had]_s []_q

action: Right-Arc(PRED)



Properties of the algorithm

- Every transition sequence outputs a projective dependency tree (**soundness**).
- Every projective dependency tree is output by some transition sequence (**completeness**).
- There are exactly $2n$ transitions in a sentence with n words.

Deterministic parsing

- If we have an **oracle** that correctly predicts the next transition $o(c)$, parsing is deterministic:

```
PARSE( $w_1, \dots, w_n$ )  
1   $c \leftarrow ([w_0]_S, [w_1, \dots, w_n]_Q, \{ \})$   
2  while  $Q_c \neq []$  or  $|S_c| > 1$   
3     $t \leftarrow o(c)$   
4     $c \leftarrow t(c)$   
5  return  $T = (\{w_0, w_1, \dots, w_n\}, A_c)$ 
```

Oracles as classifiers

- An oracle can be approximated by a (linear) **classifier**:

$$o(c) = \operatorname{argmax}_t \mathbf{w} \cdot \mathbf{f}(c, t)$$

- History-based feature representation $\mathbf{f}(c, t)$:
 - Features over input tokens relative to S and Q
 - Features over the (partial) dependency tree defined by A
 - Features over the (partial) transition sequence
- Weight vector \mathbf{w} learned from treebank data:
 - Reconstruct oracle transition sequence for each sentence
 - Construct training data set $D = \{(c, t) \mid o(c) = t\}$
 - Maximise accuracy of local predictions $o(c) = t$


Deterministic classifier-based parsing

- Advantages:
 - **Highly efficient parsing** – linear time complexity with constant time oracles and transitions
 - **Rich history-based feature representations** – no rigid constraints from inference algorithm
- Drawback:
 - Sensitive to **search errors** and **error propagation** due to deterministic parsing and local learning


Empirical results: the CoNLL 2006 shared task

- CoNLL 2006 shared task (Buchholz and Marsi 2006):
 - **MaltParser** (Nivre et al. 2006) – transition-based, deterministic, local learning
 - **MSTParser** (McDonald et al. 2006) – graph-based, exact, global learning
 - Same average parsing accuracy over 13 languages
- Comparative **error analysis** (McDonald and Nivre 2007):
 - MaltParser more accurate on short dependencies and disambiguation of core grammatical functions
 - MSTParser more accurate on long dependencies and dependencies near the root of the tree
- Hypothesised **explanation for MaltParser results**:
 - **Rich features counteracted by error propagation**

Precision by dependency length



Beam search and structured prediction



Beam search

- **Maintain the k best hypotheses** (Johansson and Nugues 2006):

```
PARSE( $w_1, \dots, w_n$ )
1  BEAM  $\leftarrow \{([w_0]_S, [w_1, \dots, w_n]_Q, \{\})\}$ 
2  while  $\exists c \in \text{BEAM} [Q_c \neq [] \text{ or } |S_c| > 1]$ 
3    foreach  $c \in \text{BEAM}$ 
4      foreach  $t$ 
5        ADD( $t(c)$ , NEWBEAM)
6    BEAM  $\leftarrow \text{TOP}(k, \text{NEWBEAM})$ 
7  return  $T = (\{w_0, w_1, \dots, w_n\}, A_{\text{TOP}(1, \text{BEAM})})$ 
```

- **Note:**

- $\text{Score}(c_0, \dots, c_m) = \sum_{i=1}^m \mathbf{w} \cdot \mathbf{f}(c_{j-1}, t_j)$
- Simple combination of locally normalised classifier scores
- Marginal gains in accuracy

Structured prediction


- **Parsing as structured prediction** (Zhang and Clark 2008):
 - Minimise loss over entire transition sequence
 - Use beam search to find highest-scoring sequence

- Factored feature representations:

$$\mathbf{f}(c_0, \dots, c_m) = \sum_{i=1}^m \mathbf{f}(c_{i-1}, t_i)$$

- Online learning from oracle transition sequences:
 - Structured perceptron (Collins 2002)
 - Early updates (Collins and Roark 2004)

Beam size and training iterations




Yue Zhang and Stephen Clark. 2008. A Tale of Two Parsers: Investigating and Combining Graph-Based and Transition-Based Dependency Parsing. In *Proceedings of the 2008 Conference on Empirical Methods in Natural Language Processing*, 562–571.

The best of two worlds?

- Like graph-based dependency parsing (MSTParser):
 - Global learning – minimise loss over entire sentence
 - Non-greedy search – accuracy increases with beam size
- Like deterministic transition-based parsing (MaltParser):
 - Highly efficient – complexity still linear for fixed beam size
 - Rich features – no constraints from parsing algorithm
- Example ZPar parser (Zhang and Clark 2011)
 - “Most heavily developed for English and Chinese”

Precision by dependency length, again




Even richer feature models

	ZPar	Malt
Baseline	92.18	89.37
+distance	+0.07	−0.14
+valency	+0.24	0.00
+unigrams	+0.40	−0.29
+third-order	+0.18	0.00
+label set	+0.07	+0.06
Extended	93.14	89.00


Yue Zhang and Joakim Nivre. 2011. Transition-Based Dependency Parsing with Rich Non-Local Features. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, 188–193.

Online reordering for non-projectivity



Projectivity

- A dependency arc is **projective** if the head (transitively) dominates all intervening words
- Most dependency grammar theories do not assume projectivity (but many parsers do)




Non-projectivity in natural languages

Language	Trees	Arcs
Arabic (Hajič et al. 2004)	11,2 %	0,4 %
Basque (Aduriz et al. 2003)	26,2 %	2,9 %
Czech (Hajič et al. 2001)	23,2 %	1,9 %
Danish (Kromann 2003)	15,6 %	1,0 %
Greek (Prokopidis et al. 2005)	20,3 %	1,1 %
Russian (Boguslavsky et al. 2000)	10,6 %	0,9 %
Slovene (Džeroski et al. 2006)	22,2 %	1,9 %
Turkish (Oflazer et al. 2003)	11,6 %	1,5 %

Projectivity and word order

- Projectivity is a property of a dependency tree only in relation to a particular word order
 - Words can always be reordered to make the tree projective
 - Given a dependency tree $T = (V, A, <)$, let the projective order $<_p$ be the order defined by an in-order traversal of T with respect to $<$ (Veselá et al. 2004)



Parsing with online reordering

- Add transition for reordering words (Nivre 2009):

- **Swap**

$$\frac{([\dots, w_i, w_j]_S, [\dots]_Q, A)}{([\dots, w_j]_S, [w_i, \dots]_Q, A)} \quad [0 < i < j]$$

- Transition-based parsing with two interleaved processes:
 - Sort words into projective order $<_p$
 - Build dependency tree T by connecting adjacent subtrees
 - T is always projective with respect to $<_p$
 - T may be non-projective with respect to $<$

Example Transition Sequence

[ROOT]_S [A, hearing, is, scheduled, on, the, issue, today, .]_Q

ROOT A hearing is scheduled on the issue today .

Example Transition Sequence

[ROOT, A]_S [hearing, is, scheduled, on, the, issue, today, .]_Q

action: Shift

ROOT A hearing is scheduled on the issue today .

Example Transition Sequence

[ROOT, A, hearing]_s [is, scheduled, on, the, issue, today, .]_q


action: Shift

ROOT A hearing is scheduled on the issue today .

Example Transition Sequence

[ROOT, A, hearing]_s [is, scheduled, on, the, issue, today, .]_q


action: Left-Arc(ATT)



Example Transition Sequence

[ROOT, hearing, is]_s [scheduled, on, the, issue, today, .]_q


action: Shift



Example Transition Sequence

[ROOT, hearing, is, scheduled]_s [on, the, issue, today, .]_q


action: Shift



Example Transition Sequence

[ROOT, hearing, is, scheduled, on]_S [the, issue, today, .]_Q


action: Shift



Example Transition Sequence

[ROOT, hearing, is, scheduled, on, the]_s [issue, today, .]_q


action: Shift



Example Transition Sequence

[ROOT, hearing, is, scheduled, on, the, issue]_s [today, .]_q


action: Shift



Example Transition Sequence

[ROOT, hearing, is, scheduled, on, the, issue]_s [today, .]_q


action: Left-Arc(ATT)



Example Transition Sequence

[ROOT, hearing, is, scheduled, on, issue]_s [today, .]_Q


action: Right-Arc(PC)



Example Transition Sequence

[ROOT, hearing, is, on]_s [scheduled, today, .]_Q


action: Swap



Example Transition Sequence

[ROOT, hearing, on]_s [is, scheduled, today, .]_Q


action: Swap



Example Transition Sequence

[ROOT, hearing, on]_s [is, scheduled, today, .]_Q


action: Right-Arc(ATT)



Example Transition Sequence

[ROOT, hearing, is]_s [scheduled, today, .]_Q


action: Shift



Example Transition Sequence

[ROOT, hearing, is]_s [scheduled, today, .]_Q


action: Left-Arc(SBJ)



Example Transition Sequence

[ROOT, is, scheduled]_s [today, .]_Q


action: Shift



Example Transition Sequence

[ROOT, is, scheduled, **today**]_s [**.**]_Q

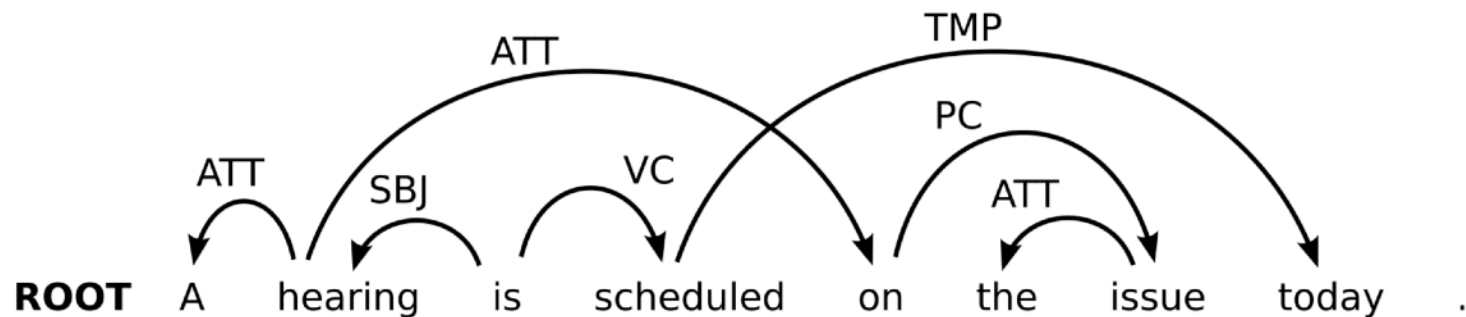
action: Right-Arc(TMP)



Example Transition Sequence

[ROOT, is, scheduled]_s [.]_Q

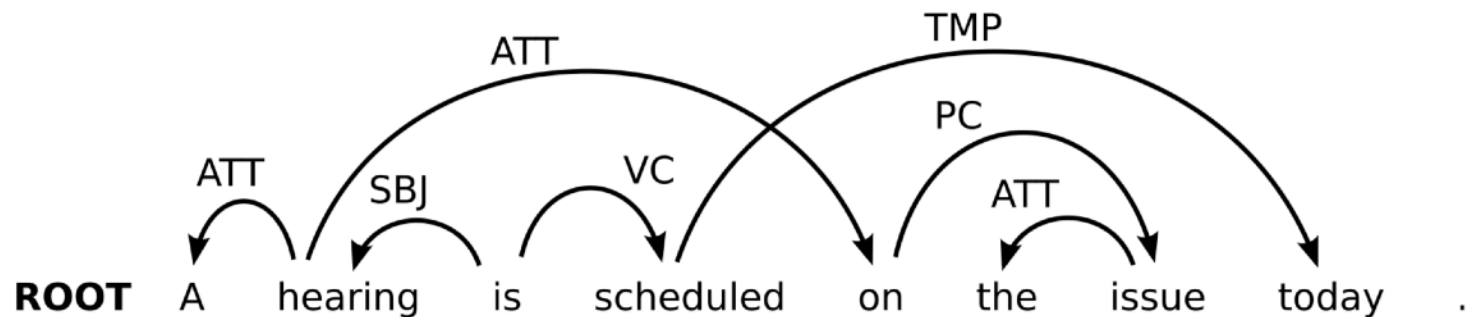
action: Right-Arc(VC)



Example Transition Sequence

[ROOT, is, .]_s []_Q

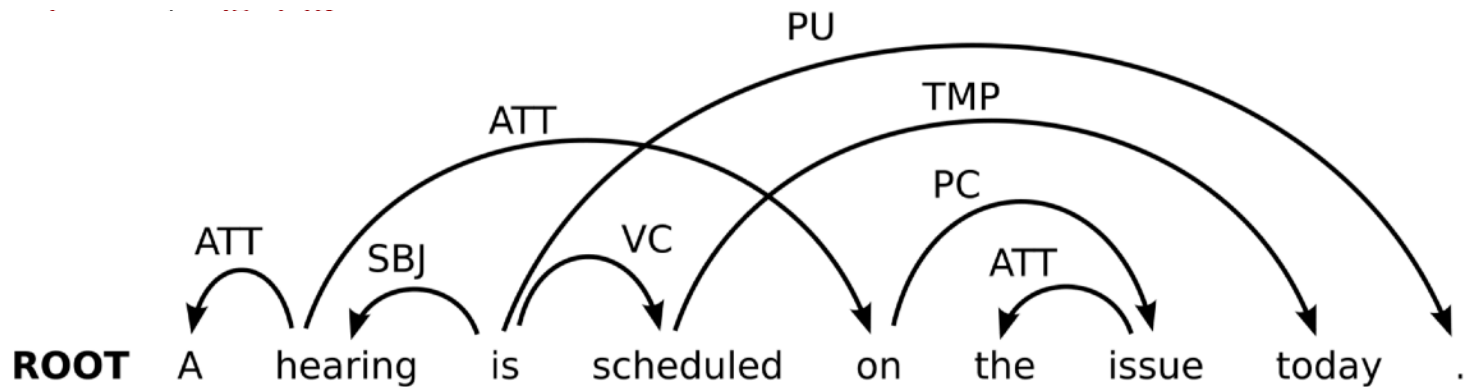
action: Shift



Example Transition Sequence

[ROOT, is, .]_s []_Q

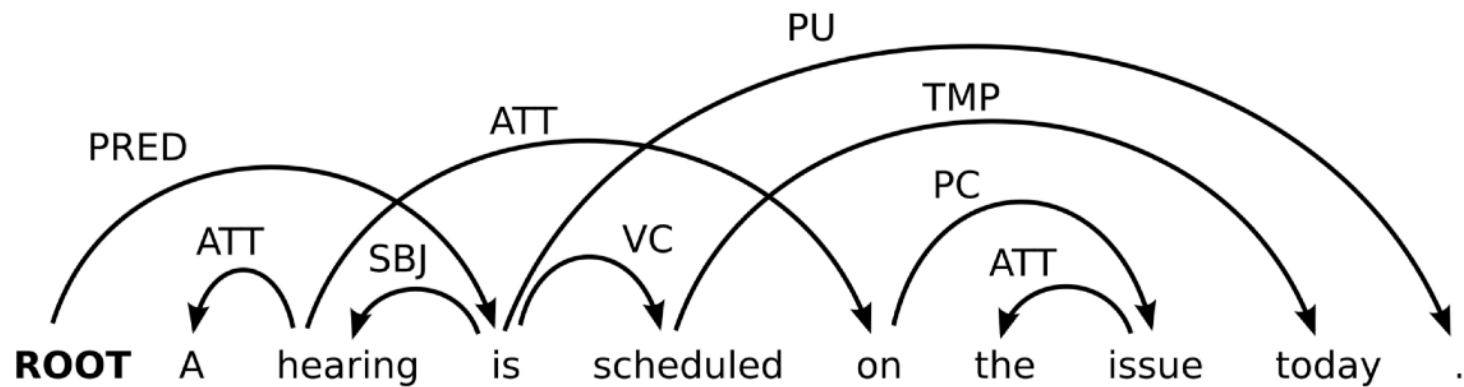
action: Right-Arc(PU)



Example Transition Sequence

[ROOT, is]_s []_q

action: Right-Arc(PRED)




Empirical results

- Deterministic transition-based parsing (Nivre 2009):
 - Parsing in linear expected time (quadratic worst-case time)
 - Best results on Czech CoNLL 2006 data sets
- Beam search and structured prediction:
 - Evaluation on CoNLL 2009 data sets (dev sets)

	Czech		German	
	LAS	UAS	LAS	UAS
Projective	80.8	86.3	86.2	88.5
Online reordering	83.9	89.1	88.7	90.9

Arc-eager Transition-Based Parsing



Limitations of the arc-standard algorithm

- The **arc-standard** system considered so far
 - builds a dependency tree strictly bottom-up
 - a dependency arc can only be added between two nodes if the dependent node has already found all its dependents.
 - As a consequence, it is often necessary to postpone the attachment of right dependents.
- This is a problem, as parsing decisions are easier to take when the governor and the governee of a dependency are immediately accessible

The problem of arc-standard on an example

$[ROOT]_S [La, \text{température}, a, \text{un}, \text{très}, \text{gros}, \text{effet}, \text{sur}, la, \text{concentration}]_Q$

La température a un très gros effet sur la concentration.

The problem of arc-standard on an example

[ROOT, La]_s [température, a, un, très, gros, effet, sur, la, concentration]_Q

action: Shift

La température a un très gros effet sur la concentration.

The problem of arc-standard on an example

[ROOT, La, température]_S [a, un, très, gros, effet, sur, la, concentration]_Q

action: Shift

La température a un très gros effet sur la concentration.

The problem of arc-standard on an example

[ROOT, La, température]_s [a, un, très, gros, effet, sur, la, concentration]_Q

action: Left-Arc()

La température a un très gros effet sur la concentration.



The problem of arc-standard on an example

[ROOT, température, a]_S [un, très, gros, effet, sur, la, concentration]_Q

action: Shift

La température a un très gros effet sur la concentration.



The problem of arc-standard on an example

[ROOT, température, a]_S [un, très, gros, effet, sur, la, concentration]_Q

action: Left-Arc()



The problem of arc-standard on an example

[ROOT, a, un]_s [très, gros, effet, sur, la, concentration]_Q

action: Shift

La température a un très gros effet sur la concentration.



The problem of arc-standard on an example

[ROOT, a, un, très]_s [gros, effet, sur, la, concentration]_q

action: Shift

La température a un très gros effet sur la concentration.

A diagram illustrating the arc-standard problem. It shows the sentence "La température a un très gros effet sur la concentration." with two curved arrows above the words. The first arrow starts above "La" and ends above "température". The second arrow starts above "température" and ends above "a".

The problem of arc-standard on an example

[ROOT, a, un, très, gros]_S [effet, sur, la, concentration]_Q

action: Shift

La température a un très gros effet sur la concentration.



The problem of arc-standard on an example

[ROOT, a, un, très, gros]_s [effet, sur, la, concentration]_Q

action: Right-Arc()



The problem of arc-standard on an example

[ROOT, a, un, gros, effet]_S [sur, la, concentration]_Q

action: Shift


La température a un très gros effet sur la concentration.



The problem of arc-standard on an example

[ROOT, a, un, gros, effet]_S [sur, la, concentration]_Q


action: Left-Arc()



The problem of arc-standard on an example

[ROOT, a, un, effet]_S [sur, la, concentration]_Q


action: Left-Arc()



The problem of arc-standard on an example

[ROOT, a, effet, sur]_s [la, concentration]_Q


action: Shift



The problem of arc-standard on an example

[ROOT, a, effet, sur, la]_s [concentration]_q


action: Shift



The problem of arc-standard on an example

[ROOT, a, effet, sur, la, concentration]_s []_q


action: Shift



The problem of arc-standard on an example

[ROOT, a, effet, sur, la, concentration]_s []_q


action: Left-Arc()



The problem of arc-standard on an example

[ROOT, a, effet, sur, concentration]_s []_Q


action: Right-Arc()



The problem of arc-standard on an example

[ROOT, a, effet, sur]_s []_Q


action: Right-Arc()



The problem of arc-standard on an example

[ROOT, a, effet]_s []_q


action: Right-Arc()



The problem of arc-standard on an example

[ROOT, a]_s []_q

action: Right-Arc()



The arc-eager system

- We will modify the basic set of actions in order to always add an arc at the earliest possible opportunity:
 - we will now build parts of the tree top-down instead of bottom-up
- Shift remains the same
- **Left-Arc is rewritten and subjected to a stricter condition**
(allowed only if the dependent is not the root and has no incoming arcs)

$$([\dots, w_i]_S, [w_j, \dots]_Q, A)$$

$$\frac{}{([\dots]_S, [w_j, \dots]_Q, A \cup \{(w_j, l, w_i)\})} \quad [i \neq 0 \wedge \nexists (k, l') \mid (k, l', i) \in A]$$

The arc-eager system

- **Right-Arc is changed:** it does not discard w_i anymore:

$$([\dots, w_i]_s, [w_j, \dots]_Q, A)$$

$$([\dots, w_i, w_j]_s, [\dots]_Q, A \cup \{(w_i, l, w_j)\})$$

- We postpone the reduction of w_i to another, new action:
- **Reduction**, only possible if the top of the stack already has a head

$$\frac{([\dots, w_i]_s, Q, A)}{([\dots]_s, Q, A)} \text{ only if } \exists(k, l') \mid (k, l', i) \in A$$

Arc-eager on an example

[ROOT]_S [La, température, a, un, très, gros, effet, sur, la, concentration]_Q

La température a un très gros effet sur la concentration.

Arc-eager on an example

[ROOT, La]_s [température, a, un, très, gros, effet, sur, la, concentration]_Q

action: Shift

La température a un très gros effet sur la concentration.

Arc-eager on an example

[ROOT, La]_s [température, a, un, très, gros, effet, sur, la, concentration]_Q

action: Left-Arc()



Arc-eager on an example

[ROOT, température]_s [a, un, très, gros, effet, sur, la, concentration]_Q

action: Shift

La température a un très gros effet sur la concentration.

A curved arrow originates from the word 'température' and points to the word 'La', indicating a dependency arc in the sentence.

Arc-eager on an example

[ROOT, température]_s [a, un, très, gros, effet, sur, la, concentration]_Q


action: Left-Arc()



Arc-eager on an example

[ROOT, a]_s [un, très, gros, effet, sur, la, concentration]_q


action: Right-Arc()



Arc-eager on an example

[ROOT, a, un]_s [très, gros, effet, sur, la, concentration]_Q


action: Shift



Arc-eager on an example

[ROOT, a, un, très]_s [gros, effet, sur, la, concentration]_q


action: Shift



Arc-eager on an example

[ROOT, a, un, très]_s [gros, effet, sur, la, concentration]_q


action: Left-Arc()



Arc-eager on an example

[ROOT, a, un, gros]_S [effet, sur, la, concentration]_Q


action: Shift



Arc-eager on an example

[ROOT, a, un, gros]_s [effet, sur, la, concentration]_Q


action: Left-Arc()



Arc-eager on an example

[ROOT, a, un]_s [effet, sur, la, concentration]_Q


action: Left-Arc()



Arc-eager on an example

[ROOT, a, effet]_s [sur, la, concentration]_Q


action: Right-Arc()



Arc-eager on an example

[ROOT, a, effet, sur]_s [la, concentration]_Q


action: Right-Arc()



Arc-eager on an example

[ROOT, a, effet, sur, la]_s [concentration]_q


action: Shift()



Arc-eager on an example

[ROOT, a, effet, sur, la]_s [concentration]_q


action: Left-Arc()



Arc-eager on an example

[ROOT, a, effet, sur, concentration]_s []_q


action: Right-Arc()



Arc-eager on an example

[ROOT, a, effet, sur, concentration]_s []_q


action: Reduce



Arc-eager on an example

[ROOT, a, effet, sur]_s []_q


action: Reduce



Arc-eager on an example

[ROOT, a, effet]_s []_q


action: Reduce



Arc-eager on an example

[ROOT, a]_s []_q


action: Reduce



Drawbacks of the arc-eager algorithm

- The arc-eager system has a weaker soundness result than the arc-standard system
- It does not guarantee the output to be a dependency tree, only a sequence of (unconnected) trees.
- In the best case, this is a sequence of length 1, meaning that the tree is in fact a tree.
- In the worst case, this is a sequence of length n , meaning that each word is its own tree.
- The arc-eager parsers normally have a last step that attaches everything that remains in the stack to the root

Transition-based parsing with a neural classifier




The problem with manual features

- Feature combinations yield literally **millions of features for parsing**
- It's very difficult to weight them all correctly or to choose the right **feature templates**
- Despite being many, they are still always incomplete
- Lexical features are extremely sparse:
 - the feature 'word surface form' can take any of **tens or hundreds of thousands categorical values...**
 - ...each absolutely unique and not related to each other
- In the end, **feature extraction sometimes takes more time than parsing itself**

Example of a neural arc-standard algorithm


- Chen and Manning (2014)
 - The first neural parsing architecture that really works
- Replace the action selection module by a neural network



Example of a neural arc-standard algorithm

- Chen and Manning (2014)
 - The first neural parsing architecture that really works
- Replace the action selection module by a neural network
- Cube activation function
 - It directly extracts feature combinations of up to three features

$$h = (W_1^w x^w + W_1^t x^t + W_1^l x^l + b_1)^3$$



Example of a neural arc-standard algorithm

- Chen and Manning (2014)
 - The first neural parsing architecture that really works
- Replace the action selection module by a neural network
- Cube activation function
- POS tags and arc labels are discrete sets
 - Normally represented as one-hot vectors
 - Just like words, there should be similarities
 - NN (singular noun) should be similar to NNP (plural noun)
 - Dense embedding layer for POS tags and arc labels capture relationships
- **Better results in accuracy and parsing speed** compared with previous parsers with statistical classifiers

Example of a neural arc-standard algorithm: Experimental Results

Parser	Dev		Test		Speed (sent/s)
	UAS	LAS	UAS	LAS	
standard	89.9	88.7	89.7	88.3	51
eager	90.3	89.2	89.9	88.6	63
Malt:sp	90.0	88.8	89.9	88.5	560
Malt:eager	90.1	88.9	90.1	88.7	535
MSTParser	92.1	90.8	92.0	90.5	12
Our parser	92.2	91.0	92.0	90.7	1013

Table 4: Accuracy and parsing speed on PTB + CoNLL dependencies.

Parser	Dev		Test		Speed (sent/s)
	UAS	LAS	UAS	LAS	
standard	90.2	87.8	89.4	87.3	26
eager	89.8	87.4	89.6	87.4	34
Malt:sp	89.8	87.2	89.3	86.9	469
Malt:eager	89.6	86.9	89.4	86.8	448
MSTParser	91.4	88.1	90.7	87.6	10
Our parser	92.0	89.7	91.8	89.6	654

Table 5: Accuracy and parsing speed on PTB + Stanford dependencies.

Parser	Dev		Test		Speed (sent/s)
	UAS	LAS	UAS	LAS	
standard	82.4	80.9	82.7	81.2	72
eager	81.1	79.7	80.3	78.7	80
Malt:sp	82.4	80.5	82.4	80.6	420
Malt:eager	81.2	79.3	80.2	78.4	393
MSTParser	84.0	82.1	83.0	81.2	6
Our parser	84.0	82.4	83.9	82.4	936

Table 6: Accuracy and parsing speed on CTB.

Example of a neural arc-standard algorithm: Model comparison





Figure 4: Effects of different parser components. Left: comparison of different activation functions. Middle: comparison of pre-trained word vectors and random initialization. Right: effects of POS and label embeddings.

Deep learning for parsing



SyntaxNet

- In 2016, Google released SyntaxNet, a neural parser implemented in TensorFlow, and state-of-the-art models:
 - <https://github.com/tensorflow/models/tree/master/research/syntaxnet>
- Implements the system described in (Andor et al. 2016):
 - ‘globally normalized transition-based dependency parser’
 - Changes compared to (Chen and Manning 2014):
 - Beam search
 - **Global optimisation** using Conditional Random Fields (CRF)
 - all valid sequences of transition operators are scored.
 - 2 hidden layers of 1024 dimensions each.
- Combines the flexibility of transition-based algorithms and the modelling power of neural networks (even without recurrence)
- Parsey McParseface model: **92.79 LAS on English PTB**
 - **LAS 80.38 on UD v1.3 English Treebank**

ParseySaurus

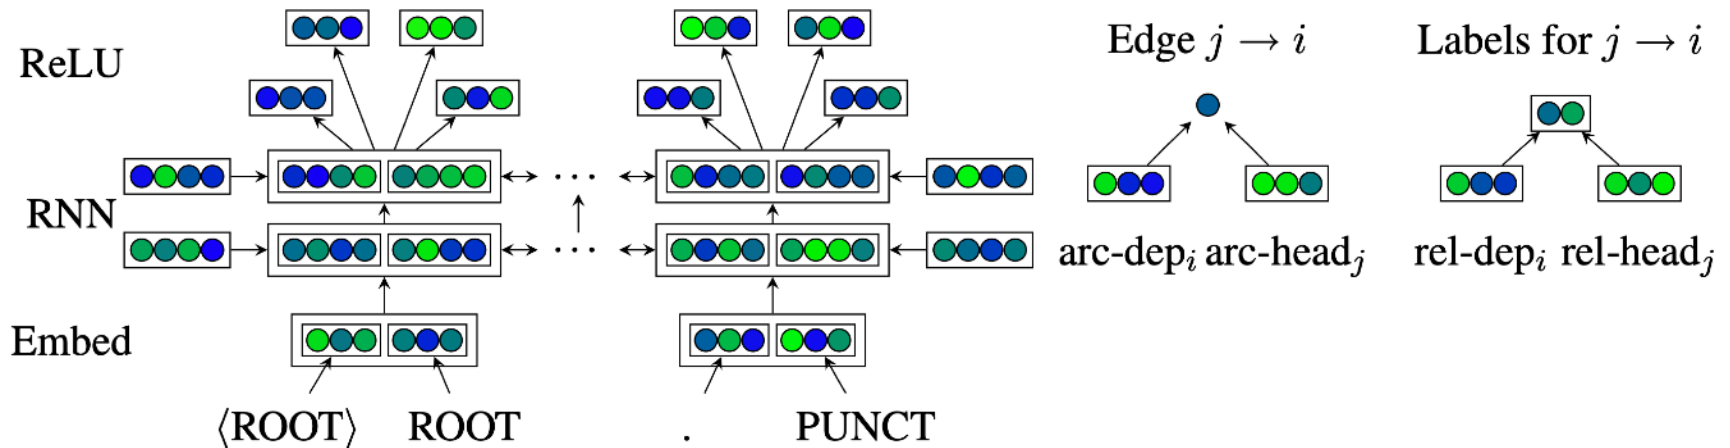
- Google then moved to using LSTMs in their **DRAGNN** framework
 - ‘Dynamic Recurrent Acyclic Graphical Neural Networks’;
 - Described in (Alberti et al. 2017)
 - LSTM transition-based neural model
 - character-based input layer
- ParseySaurus model: **84.45 LAS on UD v1.3 English Treebank**

The CoNLL 2017 shared task

- **The task was to parse raw texts in different languages into dependency trees**
- Unlike the previous CoNLL 2007 shared task, **the input is raw text**:
 - no tokenisation
 - no sentence segmentation
 - no lemmas
 - no PoS tags
- **Consistent *Universal Dependencies* (UD) annotation used for all languages**
- Training and test data came from the UD 2.0 collection:
 - 64 treebanks in 45 languages.
- 4 'surprise' languages with no training data: Buryat, Kurmanji Kurdish, North Saami and Upper Sorbian
- A major milestone in advancing data-driven dependency parsing
 - 33 participants
 - DRAGNN was one of the 2 baselines

The Dozat et al. (2017) parser

- The system described in (Dozat et al. 2017) is the winner of the shared task
 - average LAS 76.30, average UAS 81.30
- **Graph-based:** for each word, the parser looks for the most likely head, and then decides how to label the resulting dependency



The Dozat et al. (2017) parser

- The input to the model is a sequence of tokens and their part of speech tags
 - Word embeddings + character-based embeddings
- It is put through a 3-layer bidirectional LSTM network
- The output state of the final LSTM layer is then fed through four separate ReLU layers, producing four specialised vector representations for each word
 1. one for the word as a dependent seeking its head
 2. one for the word as a head seeking all its dependents
 3. another for the word as a dependent deciding on its label
 4. and a fourth for the word as head deciding on the labels of its dependents
- These vectors are then sequentially fed to two biaffine classifiers:
 - the first computes a score for each pair of tokens, with the highest score for a given token indicating that token's most probable head
 - the second computes a score for each label for a given token/head pair, with the highest score representing the most probable label for the arc from the head to the dependent

Beyond the Dozat et al. (2017) parser

- In the CoNLL 2017 shared task, Dozat and colleagues used word2vec (non-contextual) word embeddings
- They can be replaced with contextual embeddings (ELMo, BERT)
- But the contextual information provided by BERT makes the LSTM layers redundant
- The output of BERT can replace the architecture up to the LSTM layers (included)
 - This is the parsing architecture proposed by (Kondratyuk & Straka 2019)
 - It is the architecture we used to evaluate the parsing performance of our French BERT model CamemBERT (Martin et al. 2019, 2020)
 - We improve the state of the art of parsing for French



Time for a short break!